

Great—let's pin this down in math-tight, spec-ready terms using your **Engine / Verifier** split. I'll give you: formal model, constructions (incl. two non-reachability options), receipt format, verification equations, security goals, and drawing-ready figure list + captions. Minimal prose, maximal precision.

## 1) Formal model

Let:

- Directed, labeled multigraph  $G = (V, E, \lambda_V, \lambda_E)$  with
  - $E \subseteq V \times \Sigma \times V$ , edge label alphabet  $\Sigma$ .
  - $\lambda_V : V \rightarrow \{0, 1\}^*$ ,  $\lambda_E : E \rightarrow \{0, 1\}^*$  (canonical encodings).
- Policy/path spec as a **finite automaton**  $A = (Q, \Sigma, \delta, q_0, F)$  (DFA/NFA; regex compiles to DFA).
- **Relational reachability** under policy  $A$ : for  $s, t \in V$ ,

$$\text{Reach}_A(s, t) := \exists \text{path } s \xrightarrow{w} t \text{ in } G \text{ with } w \in L(A).$$

Define **product graph**  $H = V \times Q$  with edges

$$(u, q) \rightarrow (v, q') \iff \exists a \in \Sigma \text{ s.t. } (u, a, v) \in E \wedge q' = \delta(q, a).$$

Then  $\text{Reach}_A(s, t)$  iff  $\exists q_f \in F$  with  $(t, q_f) \in \text{Post}^*(\{(s, q_0)\}, H)$ .

## 2) Commitments and hashing

Let  $H(\cdot)$  be a collision-resistant hash (e.g., SHA-256). Canonical element encodings:

- $\langle v \rangle := \text{CBOR}(\lambda_V(v))$ ,
- $\langle e \rangle := \text{CBOR}(u, a, v, \lambda_E(e))$ ,
- $\langle u, a, v \rangle$  canonical triple.

### 2.1 Dataset commitment (Engine)

Build a key-sorted **sparse Merkle map** (SMM) over:

- Node set  $V$ : keys  $k_V = H("V" \parallel \langle v \rangle)$ , values 1.
- Edge set  $E$ : keys  $k_E = H("E" \parallel \langle u, a, v \rangle)$ , values  $H(\lambda_E(e))$ .

Let  $R_V, R_E$  be the roots. Define dataset root:

$$R_D = H("D" \parallel R_V \parallel R_E).$$

### 2.2 Policy commitment

$$R_A = H("A" \parallel \text{CBOR}(Q, \Sigma, \delta, q_0, F)).$$

### 3) Receipts

A receipt proves statement  $S = (s, t, A, \text{mode})$  at time  $\tau$  relative to dataset commitment  $R_D$ .

#### 3.1 Receipt structure (common)

cpp

 Copy code

```
Receipt = {
    version,                                // e.g., 1
    statement: {s_id, t_id, A_id, mode},      // ids or encodings
    RD, RA,                                    // commitments
    anchor: TSA_proof,                      // RFC 3161 blob or dual anchors
    witness: W,                            // differs by mode
    sig: Sig_SK( H( canonicalize(all_above) ) )
}
```

Engine signs with keypair  $(\text{SK}, \text{PK})$ ; Verifier knows  $\text{PK}$ .

#### 3.2 Reachability witness $W^{\text{reach}}$

A path witness:

- Node/edge sequence  $(v_0 = s, a_1, v_1, \dots, a_m, v_m = t)$ .
- For each used edge  $(v_{i-1}, a_i, v_i)$ : SMM inclusion proof  $\pi_E^i$  w.r.t. root  $R_E$ .
- DFA trace  $q_0 \xrightarrow{a_1} q_1 \cdots \xrightarrow{a_m} q_m$  with  $q_m \in F$ .

Verifier checks:

1.  $\forall i, \text{SMM.VerifyInclude}(R_E, k_E^i, H(\lambda_E(e_i)), \pi_E^i) = \text{true}$ .
2.  $q_i = \delta(q_{i-1}, a_i)$  and  $q_m \in F$ .
3. Signature and anchor validate.



### 3.3 Non-reachability witnesses $W^{\text{non-reach}}$ (two constructions)

#### (A) Product-state sparse map (PSSM) + cut certificate (practical, patentable)

Engine computes **closure bitmap**  $B : V \times Q \rightarrow \{0, 1\}$  for reachable product-states starting from  $(s, q_0)$ .

Builds an SMM over keys

$k_H = H(\text{"H"} \parallel \langle v \rangle \parallel q)$  with value  $B(v, q)$ . Root  $R_H$ .

To avoid “trust me”, include a **local cut certificate**  $C$  consisting of a minimal **frontier set**  $F^* \subseteq V \times Q$  such that:

- Every path from  $(s, q_0)$  to any  $(t, q_f)$  must pass through  $F^*$ .
- All  $(x, q) \in F^*$  are **closed**: for every outgoing transition  $(x, q) \rightarrow (y, q')$  permitted by  $E$  and  $\delta$ , we have  $B(y, q') = 0$ .

**Witness**  $W^{\text{non-reach}} = (R_H, \Pi_0, C, \Pi_C)$ :

- $\Pi_0$ : SMM non-inclusion proof for the target states: for each  $q_f \in F$ ,  
 $\text{SMM.VerifyZero}(R_H, k_H(t, q_f), \pi_{t, q_f})$ .
- $C = \{(x_j, q_j)\}_{j=1}^r$  with:
  - ◆ Inclusion proofs  $\pi_{x_j, q_j}$  showing  $B(x_j, q_j) = 1$ .
  - ◆ For each outgoing edge  $(x_j, a, y)$  with  $q' = \delta(q_j, a)$ : proofs  $\pi_{y, q'}$  showing  $B(y, q') = 0$  (or alternately, proofs of edge absence from  $R_E$  if policy forbids it).
- $\Pi_C$ : SMM proofs supporting the above claims.

**Verifier checks**:

- ◆ Non-inclusion at  $(t, q_f)$  for all  $q_f \in F$ .
- ◆ Each frontier state is marked 1; each of its policy-consistent successors is 0 (or edge absent).

By **frontier minimality**, no accepting target is reachable  $\Rightarrow$  non-reachability holds.

*Note:* Frontier minimality can be encoded as: each  $(x_j, q_j)$  has at least one predecessor in  $B = 1$  (prove inclusion) and all successors are 0; the union of these cuts all  $s \rightarrow F$  paths (Engine ensures; Verifier checks the local conditions).

#### (B) Edge-cut certificate over $G$ (small-footprint variant)

Provide a **cut set of edges**  $K \subseteq E$  such that every  $A$ -conforming  $s \rightarrow t$  path uses some edge in  $K$ .

Witness includes:

- Inclusion proofs that all edges in  $K$  exist in  $R_E$ .
- Proofs (policy side) that removing  $K$  destroys all automaton-conforming paths (encode each edge's label position in the automaton; Verifier simulates supergraph with edges in  $K$  removed to confirm no residual accepting path exists using small “blocking lemmas” tied to  $A$ ).

This variant trades more policy algebra for fewer SMM proofs; include if you want a second embodiment.

#### 4) Verification equations (succinct)

- Dataset root:  $R_D \stackrel{?}{=} H("D" \parallel R_V \parallel R_E)$ .
- Policy root:  $R_A \stackrel{?}{=} H("A" \parallel \text{CBOR}(A))$ .
- Anchor:  $\text{RFC3161.Verify}(R_D \parallel R_A, \text{TSA\_proof}) = \text{true}$ .
- Signature:  $\text{VRF} = H(\text{canonical Receipt minus sig})$ ;  $\text{Sig.Verify}(PK, \text{VRF}, \text{sig}) = \text{true}$ .
- Reachability:  $\bigwedge_i \text{SMM.Incl}(R_E, k_E^i, H(\lambda_E(e_i)), \pi_E^i)$  and DFA trace acceptance.
- Non-reachability (PSSM):

$$[\forall q_f \in F : \text{SMM.Zero}(R_H, k_H(t, q_f), \pi)] \wedge \bigwedge_{(x,q) \in C} (\text{SMM.Incl}(R_H, k_H(x, q), 1, \pi) \wedge \forall (x, q) \rightarrow (y, q') : \text{SMM.Zero}(R_H, k_H(y, q'), \pi)).$$

## 5) Security goals

Assume collision-resistant  $H$ , EUF-CMA signature, and sound RFC-3161.

♦ **Receipt Unforgeability:** No PPT adversary can produce a valid Receipt asserting a false statement  $S$  w.r.t.  $(R_D, R_A)$  unless it:

- finds a Merkle collision (breaks CRH),
- forges signature,
- forges TSA proof.

♦ **Soundness (reachability):** A false positive requires including a non-existent edge (Merkle collision) or mis-simulating  $A$ .

♦ **Soundness (non-reachability):** A false non-reach receipt requires either:

- marking an actually-reachable product-state as 0 (Merkle collision), or
- presenting a frontier  $C$  where some successor is actually 1 (contradicted by inclusion proofs), or
- eliding an existing successor edge (contradicted by  $R_E$  inclusion).

♦ **Completeness:** If statement is true, Engine can construct witnesses (path or valid frontier/cut) in  $O(m \log n)$ .

## 6) Complexity

Let  $n = |V|, m = |E|, |Q| = q$ .

- ♦ Path witness size:  $O(\ell \cdot \log N)$  hashes, where  $\ell$  is path length and  $N$  SMM size. Verify time  $O(\ell \log N)$ .
- ♦ Non-reach (PSSM): frontier size  $r$ , out-degree  $\Delta$ : witness includes  $r$  inclusions +  $r\Delta$  zero-proofs  $\Rightarrow$  size  $O((r + r\Delta) \log N)$ . In many RBAC/data-lineage graphs,  $r \ll n$ .
- ♦ Microsecond-level verification on commodity CPUs ( 20 SHA-256 for  $\log_2 N \approx 20$ ).

## 7) Engine / Verifier separation (math view)

### Engine

- ◆ Builds  $R_V, R_E, R_D$ ; compiles  $A \rightarrow R_A$ .
- ◆ Reachability: extracts path  $\pi$  + inclusion proofs.
- ◆ Non-reachability: computes closure  $B$  over  $H$ ; derives minimal frontier  $C$ ; constructs  $R_H$  and the proof set  $\Pi$ .
- ◆ Produces Receipt and  $\text{Sig}_{\text{SK}}$ .

### Verifier

- ◆ Recomputes no closures; it **only**:

- checks Merkle (inclusion/zero) proofs against roots,
- checks DFA steps on provided symbols,
- checks TSA + signature.

## 8) Pseudocode (spec-ready)

pseudo

 Copy code

```
function Engine.GenerateReceipt(G, A, s, t, mode, TSA, SK):
    (RV, RE) := SMM.BuildNodesEdges(G)
    RD := H("D" || RV || RE)
    RA := H("A" || CBOR(A))
    if mode == "reach":
        path := FindPath_A(G, A, s, t)           // any A-conforming path
        W := []
        for each edge e=(u,a,v) in path:
            k := H("E" || <u,a,v>)
            W.append( SMM.ProveInclude(RE, k, H(lambdaE(e))) )
        witness := {path_labels, W}
    else: // "nreach"
        B := ClosureBitmap(G, A, s)           // over product graph H
        (RH, MapH) := SMM.BuildBitmap(B)
        C := Frontier(B, A)                  // minimal frontier cut
        witness := {RH,
                    ZeroProofsForTargets(MapH, t, A.F),
                    ProofsForFrontier(MapH, C),
                    ZeroProofsForSuccessors(MapH, C, G, A) }
    anchor := TSA.Timestamp( RD || RA )
    payload := canonicalize(version, stmt, RD, RA, anchor, witness)
    sig := Sign(SK, H(payload))
    return Receipt(payload, sig)
```



pseudo

Copy code

```
function Verifier.Verify(Receipt, PK):
    parse (version, stmt, RD, RA, anchor, witness, sig)
    assert SigVerify(PK, H(canonicalize(...)), sig)
    assert RFC3161.Verify(RD || RA, anchor)
    if stmt.mode == "reach":
        for each edge proof in witness:
            assert SMM.VerifyInclude(RE_from_RD, key, val, proof)
        assert DFA_Accepts(stmt.A, witness.path_labels)
        return True
    else: // "nreach"
        assert forall qf in A.F:
            SMM.VerifyZero(RH_from_witness, k_H(t,qf), π)
        for each (x,q) in C:
            assert SMM.VerifyInclude(RH, k_H(x,q), 1, π)
            for each policy-consistent successor (y,q'):
                assert SMM.VerifyZero(RH, k_H(y,q'), π)
        return True
```

## 9) Figures (patent drawing set)

- ◆ FIG. 1 System context: data sources → Engine → Receipt → Verifier/Auditor.
- ◆ FIG. 2 Dataset commitment: (a) node SMM  $R_V$ ; (b) edge SMM  $R_E$ ; (c) combined  $R_D$ .
- ◆ FIG. 3 Reachability receipt: path witness with per-edge Merkle proofs into  $R_E$ ; DFA trace overlay.
- ◆ FIG. 4 Product graph  $H = V \times Q$ : closure bitmap  $B$ ; minimal frontier  $C$ ; SMM  $R_H$ .
- ◆ FIG. 5 Non-reachability verification flow: (i) zero-proofs at  $(t, q_f)$ ; (ii) frontier inclusions; (iii) successor zero-proofs.
- ◆ FIG. 6 Anchoring timeline:  $R_D \| R_A \rightarrow$  TSA stamp; receipt signing; audit timeline.
- ◆ FIG. 7 Alternative embodiment: edge-cut certificate  $K$  with inclusion proofs and automaton blocking lemmas.
- ◆ FIG. 8 Receipt structure exploded view (fields, sizes, canonicalization order).

## **10) Claims-adjacent phrasing (for your spec)**

- "A method for generating a cryptographic receipt certifying relational non-reachability, comprising: committing to a dataset with a collision-resistant Merkle map; compiling a path policy to a finite automaton; constructing a product-state closure over a graph of nodes and automaton states; deriving a frontier set that intercepts all policy-conforming paths; and producing Merkle inclusion and non-inclusion proofs sufficient for independent verification without recomputing the closure."
- "A verifier configured to validate said receipt by checking only commitment-consistent Merkle proofs, automaton transitions on provided labels, and an external timestamp, wherein no access to the underlying dataset is required."